

Dhruv Patel

Los Angeles, CA | U.S. Citizen | (714) 707-1591 | dhruvpatel97@g.ucla.edu | [LinkedIn](#) | [GitHub](#) | [Website](#)

EDUCATION

University of California, Los Angeles (UCLA)

Jun. 2027

B.S. – Mathematics of Computation | Minor – Data Science Engineering

GPA: 3.77/4.00

Coursework: Machine Learning, Deep Learning & Neural Networks, ML for Finance (Graduate), AI for Scientific Computing (Graduate), NLP, Real Analysis, Probability, Linear Algebra, Data Structures, Algorithms, Numerical Methods, Theory of Computing

Honors: Upsilon Pi Epsilon (CS Honor Society), UCLA Dean's Honors List (3×), Eagle Scout

Organizations: Association for Computing Machinery, UCLA Poker Group, UCLA Cubing Club, Intramural Basketball

EXPERIENCE

UCLA Mathematical & Statistical Finance Group

Los Angeles, CA

Quantitative Research Assistant

Jun. 2025 – Present

- Researching ML and AI applications in quantitative finance under Prof. Mihai Cucuringu
- Forecasted daily volume dispersion over a ~4k stock universe; HAR-based model cut OOS forecast error 18% vs AR(1) baseline
- Built market-neutral long-short equity strategy using volume dispersion; 2.81 Sharpe over 25-year backtest (2000–2024)
- Designed configurable strategy backtesting pipeline, reaching 4× speedup via vectorized NumPy operations and 5D PnL tensor

Pacific Life

Newport Beach, CA

Data Science Intern

Jun. 2026 – Sep. 2026

- Implemented a Cortex Search Service to filter medical records down to only the information relevant to underwriting guidelines
- Compared filtered records to applicant answers via LLM, projected to cut caseload ~10% by skipping review of no-discrepancy cases
- Looped an agent through the detection system, tuning a YAML search space until false omission rate hits target (5-loop max)
- Developed an ADO pipeline pulling per-user Snowflake spend and triggering AWS SNS email alerts on threshold breaches

Smart Circle International

Newport Beach, CA

AI Engineer Intern

Sep. 2025 – Jun. 2026

- Developed a Snowflake Cortex LLM agent for D2D sales territory reporting and zip code-level lead replacement recommendations
- Built a Python pipeline to parse and structure SCORM e-learning content across 10+ formats for downstream AI training workflows
- Aggregated 15+ campaign documents across formats and structures into a centralized knowledge base for LLM ingestion

Analytics Intern

Jun. 2025 – Sep. 2025

- Built Python churn prediction models (Random Forest & Logistic Regression), achieved 50%+ recall with 85%+ accuracy
- Synthesized 100K+ dispositions into a weekly Tableau report on sales efficiency by office; adopted across 15+ sales campaigns
- Designed office-level metrics evaluating AI-driven recruiting; informed funding allocation decisions across 100+ offices

Bruin Sports Analytics

Los Angeles, CA

Tennis Data & Scouting Analyst

Oct. 2024 – Jun. 2026

- Built a transformer model to predict shot type and court location from sequential rally data using PyTorch
- Created 20+ visualizations of serving tendencies used by the UCLA D1 Men's Tennis Team to adjust pre-match strategy
- Standardized end-to-end workflow from match footage to court visualizations, enabling adoption across the scouting team

PROJECTS

LOBSTER News Alpha Agent | *Python (pandas, SciPy, NumPy), Anthropic API*

- Built an LLM agent combining news sentiment with limit-order-book order flow for intraday signals on 10 NASDAQ names
- Controlled for data contamination via anonymization placebo test; applied BH-FDR and bootstrap CIs across 2,879 events
- Identified one subgroup surviving FDR: agent beat lexicon baseline on inference-heavy TSLA headlines (IC +0.28, $p < 0.001$)

NBA Timeout Optimization via Hidden Markov Models | *Python (hmmlearn, NumPy, pandas, Scikit-Learn)*

- Fit a 3-state HMM on 250K+ NBA possessions to infer latent momentum regimes from sequential scoring data
- Applied within-team covariate matching across 12K+ pairs to estimate timeout effects net of team quality
- Found timeout impact peaks in Q4 when trailing 6-10 points (+0.14 points/possession)

EEG Motor Imagery Classification | *Python (PyTorch, NumPy, SciPy)*

- Built a data pipeline to resample, bandpass filter, and normalize noisy 22-channel EEG signals across 9 subjects
- Benchmarked five architectures (CNN, LSTM, GRU, TCN, Transformer) under a 36-fold leave-one-subject-out protocol
- Showed compact models (4K parameters) generalize better than 200K-parameter alternatives

SKILLS

Languages: Python, C++, Java, C, SQL, R, JavaScript, HTML/CSS, LaTeX

Libraries: NumPy, pandas, Scikit-Learn, PyTorch, SciPy, hmmlearn, Optuna, LightGBM, XGBoost, Matplotlib, Seaborn

Tools: Snowflake, Azure DevOps, Docker, Git, Linux, JupyterLab, Tableau, Excel, Google Cloud